

02-Creating an Embedding



1- Go to [platform.openai.com](https://platform.openai.com/docs/api-reference/embeddings) and go on api reference tab and select embedding (<https://platform.openai.com/docs/api-reference/embeddings>)

OpenAI Platform

Dashboard Docs API reference

Search (x) Introducing GPT-4o: our fastest and most affordable flagship model

GET STARTED

- Overview
- Introduction
- Quickstart
- Models
- Tutorials
- Prompt examples
- Changelog

CAPABILITIES

- Text generation
- Function calling
- Embeddings**
- Overview
- Models
- Use cases
- Frequently asked questions
- Fine-tuning
- Image generation
- Vision

Forum Help

Embeddings

Learn how to turn text into numbers, unlocking use cases like search.

New embedding models

text-embedding-3-small and text-embedding-3-large, our newest and most performant embedding models are now available, with lower costs, higher multilingual performance, and new parameters to control the overall size.

What are embeddings?

OpenAI's text embeddings measure the relatedness of text strings. Embeddings are commonly used for:

- **Search** (where results are ranked by relevance to a query string)
- **Clustering** (where text strings are grouped by similarity)
- **Recommendations** (where items with related text strings are recommended)
- **Anomaly detection** (where outliers with little relatedness are identified)
- **Diversity measurement** (where similarity distributions are analyzed)
- **Classification** (where text strings are classified by their most similar label)

An embedding is a vector (list) of floating point numbers. The [distance](#) between two vectors measures their relatedness. Small distances suggest high relatedness and large distances suggest low relatedness.

Visit our [pricing page](#) to learn about Embeddings pricing. Requests are billed based on the number of [tokens](#) in the input.

How to get embeddings

2- Create the postman request and add the secret Key by fetching it from openAI.

HTTP Open AI vector DB / Create Embeddings

Save

POST https://api.openai.com/v1/embeddings

Send

Params Authorization Headers (8) Body Pre-request Script Tests Settings

Get your API Key for OpenAI API

- Go to [OpenAI API Keys page](#)
- Click on **Create new secret key**.
- Copy the key and paste it here.

API Key

Visit platform.openai.com/account/api-keys to get your API Key.

Enter the Bearer Token here

[Save authorization to collection](#)

Create secret key from openAI

DashboardDocsAPI reference

API keys

Create new secret key

Project API keys have replaced user API keys.

We recommend using project based API keys for more granular control over your resources. [Learn more](#)

View user API keys

As an owner of this project, you can view and manage all API keys in this project.

Do not share your API key with others, or expose it in the browser or other client-side code. In order to protect the security of your account, OpenAI may also automatically disable any API key that has leaked publicly.

View usage per API key on the [Usage page](#).

Create an API key to access the OpenAI API.

Create new secret key

POST Create Embeddings

Open AI vector DB / Create Embeddings

POSThttps://api.openai.com/v1/embeddings

Send

ParamsAuthorizationHeaders (9)BodyPre-request ScriptTestsSettingsCookies

Need help authorizing with OpenAI API?

TypeBearer Token

The authorization header will be automatically generated when you send the request. Learn more about [Bearer Token](#) authorization.

Heads up! These parameters hold sensitive data. To keep this data secure while working in a collaborative environment, we recommend using variables. [Learn more about variables.](#)

Token

sk-proj-2plb3b4e9dg6Z94J3bPJT3BlbkFJr...

Response

Looking at the documentation, we only require model and input param as necessary for the api.

Create embeddings

POST https://api.openai.com/v1/embeddings

Creates an embedding vector representing the input text.

Request body

model string Required

ID of the model to use. You can use the [List models](#) API to see all of your available models, or see our [Model overview](#) for descriptions of them.

input string or array Required

Input text to embed, encoded as a string or array of tokens. To embed multiple inputs in a single request, pass an array of strings or array of token arrays. Each input must not exceed the max input tokens for the model (8191 tokens for `text-embedding-ada-002`). [Example](#) [Python code](#) for counting tokens.

user string Optional

A unique identifier representing your end-user, which can help OpenAI to monitor and detect abuse. [Learn more](#).

```
Example requestcurlhttps://api.openai.com/v1/embeddings \
-H "Authorization: Bearer $OPENAI_API_KEY" \
-H "Content-Type: application/json" \
-d '{
  "input": "The food was delicious and the waiter..",
  "model": "text-embedding-ada-002"
}'
```

```
ParametersCopy{
  "model": "text-embedding-ada-002",
  "input": "The food was delicious and the waiter..."
}
```

```
ResponseCopy{
  "object": "list",
  "data": [
    {
      "object": "embedding",
      "embedding": [

```

Ensure you have credit in your account

3- create body request and send to get the embedding.

to Embeddings

c:/api.openai.com/v1/embeddings

to Embeddings

Save

POST

https://api.openai.com/v1/embeddings

Send

Params

Authorization

Headers (32)

Body

Pre-request Script

Tests

Settings

none

form-data

x-www-form-urlencoded

raw

binary

GraphQL

JSON

Beautify

1

{

2

... "model": "text-embedding-ada-002",

3

... "input": "Hello World"

4

}

Body

Cookies (24)

Headers (24)

Test Results

Pretty

Raw

Preview

Visualize

[{"object": "list", "data": [{"object": "embedding", "index": 0, "embedding": [-0.007072272, 0.0035393215, -0.007135986, -0.02920657, -0.01304866, 0.0109970635, -0.020273846, 0.0053201322, -0.008569554, -0.030251483, 0.024377037, 0.0098056095, -0.027499031, -0.0065689296, 0.009130239, 0.01548254, 0.021892184, -0.009741895, 0.011054407, 0.015890311, -0.0057756887, 0.0070340433, 0.00531321755, 0.004769005, -0.014017115, -0.0063363733, 0.005520832, -0.025434691, 0.03091411, -0.03027697, 0.009811981, -0.012188518, -0.0024625522, -0.02803423, 0.00779499, -0.0269186, -0.0026186518, -0.014047342, 0.018808419, -0.01939459, 0.0049824473, 0.014221005, 0.007524642, -0.019458305, -0.036801297, 0.015176713, 0.003533393, -0.010920979, -0.0036586213, 0.023740521, 0.032993336, 0.004769605, -0.015112999, 0.007397214, -0.0116342055, 0.0038132924, -0.0238416514, 0.0238402972, 0.027014803, -0.025001436, 0.007537385, 0.00958261, -0.0056673745, 0.0047817477, -0.00040478402, 0.011137235, 0.013125117, 0.0071168714, -0.004373977, 0.022796927, 0.026734462, 0.0066135298, 0.0043644197, -0.010141294, 0.021662815, 0.00596046, -0.015469798, -0.0015697576, 0.006263102, 0.0048295334, 0.033844963, -0.027193204, -0.009385096, 0.008034355, 0.019866075, 0.004179649, 0.002253884, 0.0143994, -0.009213068, -0.020617802, 0.008849865, 0.014157286, 0.004944219, 0.010117808, -0.01007958, 0.0038132924, -0.010210751, 0.013278031, 0.0021869885, -0.04490574, -0.013239802, 0.00429752, -0.009041039, -0.014157286, -0.0069193575, -0.001712318, -0.0041318634, 0.002792273, 0.027278403, 0.007429071, -0.016527453, 0.012832032, -0.016030483, -0.036724843, 0.0029945655, -0.007263414, 0.0039853207, -0.007945156, -0.0217775, -0.022108814, 0.014972827, 0.040394776, 0.019853331, -0.022376413, 0.007690299, 0.007798613, -0.043605972, -0.014602485, -0.004488624, -0.00909182, 0.039579235, 0.002562902, 0.011513148, -0.0098056095, -0.028747829, 0.015813854, -0.010544693, 0.021994127, -0.013290773, -0.015866426, 0.0021216816, 0.021701043, 0.013889687, -0.011271034, -0.00497289, 0.0026457305, 0.006893872, -0.01088875, 0.0029404084, -0.015609969, 0.0021137171, 0.016183397, 0.0037209068, -0.00794728, 0.011908176, 0.01619614, -0.002628209, 0.023816353, 0.0055813603, -0.007186957, 0.00036496267, -0.0003719314, -0.0019576175, -0.00355525, 0.011787119, 0.026607033, 0.027728403, -0.0066412643, -0.009907352, -0.02765583, -0.0104427505, 0.013074146, -0.030302454, 0.018744705, -0.014348429, -0.005619589, -0.0026712161, 0.023000812, -0.010474407, -0.018961335, -0.02286125, 0.010920607, -0.002542195, 0.016208882, 0.006365294, -0.0019846961, 0.011595148, -0.00886264, 0.012360547, -0.020554189, 0.018910363, 0.03014954, 0.030861338, 0.020919888, -0.6936178, -0.007671185, 0.0128702605, 0.0067600724, 0.016387282, 0.005281904, 0.014450371, 0.025778748, 0.0062089446, 0.022083327, -0.010633893, 0.017801736, -0.008270998, -0.0059286025, 0.00032713238, -0.008856268, -0.0059923166, -0.026148291, -0.023574239, 0.022134298, 0.0080330623, 0.025766006, -0.02394753, -0.013498544, 0.001865232, 0.00907552, 0.0067919292, -0.0032685364, -0.017292023, 0.028187145, -0.0078049847, 0.019203447, 0.0009294303, 0.0076265847, 0.067791864, 0.020261103, -0.012182147, 0.012647261, 0.0052086324, 0.030939596, -0.015431569, -0.015087513, 0.0036858642, -0.01660391, -0.005323318, 0.0018238178, 0.010570179, -0.006925729, 0.00030284136, 0.005508089, 0.027116746, -0.012832032, 0.01831145, 0.012022862, 0.00020378576, 0.0024020239, 0.01989156, 0.01066575, -0.0039566495, 0.0009260696, -0.0007462521, -0.016438253, -0.0067278152, -0.0010815478, -0.018846648, 0.028722342, -0.02380361, 0.0013356081, -0.0014813542, -0.0057215313, -0.0018700105, -0.00336136, -0.022665676, -0.024262352, 0.01566094, 0.033972739, 0.00823284, -0.012679118, -0.005036604, 0.024066606, 0.013927915, 0.0050684614, -0.02047773, -0.018290706, 0.01223949, -0.017292023, -0.01366229, -0.014348429, -0.000684529, 0.0056673745, 0.030175027, 0.007792242, -0.013558373, -0.018591791, 0.026989318, -0.004042636, -0.01665488, 0.006084877, 0.07447898, -0.00075669564, -0.00046157048, 0.0011165007, 0.0077071168, 0.011031667, 0.075586081, -0.01627451, -0.0057864711, 0.071306016, 0.077760457, -0.01608076

you can send 1 word, 2 words or large chunks of paragraph and get the embeddings. (why are they almost of the same size?)
These embeddings can then be used to search in database.

"text-embedding-ada-002" model takes 8192 token which is approx around (8192*4= 32764 letters) (around 10 pages of pdf approx)